

Azure RAG Architecture for SharePoint

Custom RAG Solution for Enterprise Environments

Date: February 2026 **Purpose:** Design document for building a custom RAG (Retrieval-Augmented Generation) system using Azure services to query SharePoint documents.

Table of Contents

1. [Overview](#)
 2. [Architecture Diagram](#)
 3. [Components](#)
 4. [Document Ingestion Pipeline](#)
 5. [Query Pipeline](#)
 6. [Azure Services Setup](#)
 7. [Code Examples](#)
 8. [Cost Estimates](#)
 9. [Comparison with M365 Copilot](#)
 10. [UI Options](#)
 11. [Implementation Roadmap](#)
-

Overview

This document outlines a custom RAG solution for organizations that:

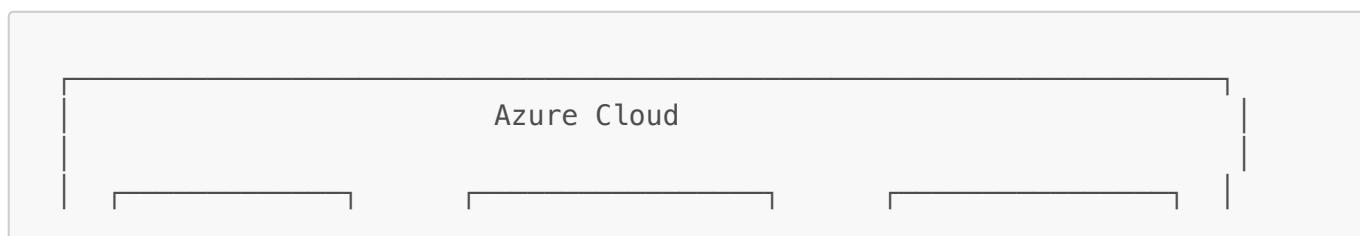
- Use Azure as their enterprise cloud provider
- Store documents in SharePoint Online
- Cannot use external AI services (e.g., Claude/Anthropic)
- Want to avoid M365 Copilot licensing costs (\$30/user/month)

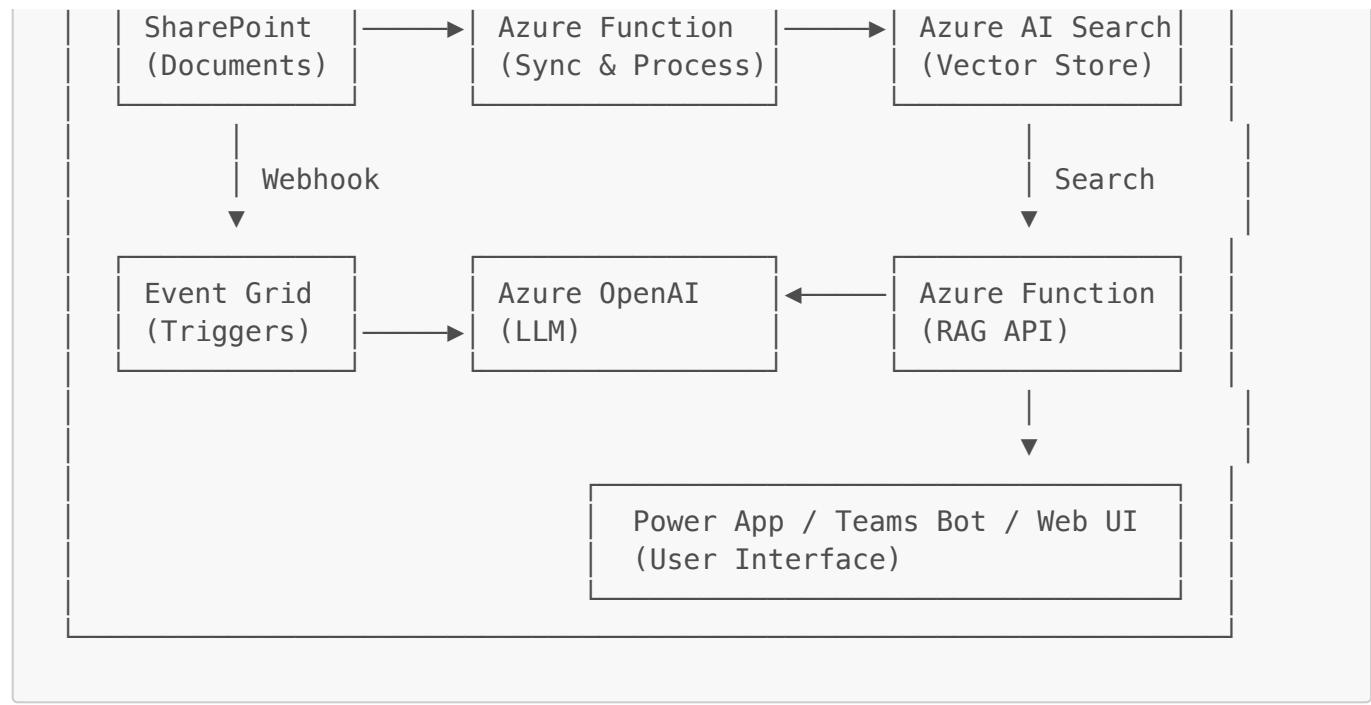
What is RAG?

Retrieval-Augmented Generation combines:

1. **Retrieval:** Find relevant documents using semantic search
 2. **Augmented:** Add retrieved content to the AI prompt
 3. **Generation:** AI generates answers based on your documents
-

Architecture Diagram





Components

Component	Azure Service	Purpose
Documents	SharePoint Online	Source of truth for all documents
Change Detection	Event Grid + Webhooks	Detect when files are added/modified
Processing	Azure Functions	Parse documents, chunk text, generate embeddings
Vector Database	Azure AI Search	Store embeddings and perform semantic search
LLM	Azure OpenAI (GPT-4)	Generate natural language answers
User Interface	Power Apps / Teams / Web App	End-user interaction
Authentication	Azure AD / Entra ID	Enterprise SSO and access control

Document Ingestion Pipeline

When a document is uploaded or modified in SharePoint:



- 1. Download file from SharePoint (via Graph API)
- 2. Parse content (PDF, Word, Excel, PowerPoint)
- 3. Chunk text into segments (500–800 characters)
- 4. Generate embeddings (Azure OpenAI text-embedding-ada-002)
- 5. Store chunks + embeddings in Azure AI Search

Supported File Formats

Format	Extension	Parser
PDF	.pdf	PyMuPDF / Azure Document Intelligence
Word	.docx	python-docx
Excel	.xlsx	openpyxl
PowerPoint	.pptx	python-pptx
Text	.txt, .md	Direct read

Query Pipeline

When a user asks a question:

- ```

User asks question
 ↓
Azure Function: "QueryRAG"
 └── 1. Convert question to embedding vector
 └── 2. Search Azure AI Search (vector similarity)
 └── 3. Retrieve top 5 relevant document chunks
 └── 4. Build prompt with context + question
 └── 5. Send to Azure OpenAI (GPT-4)
 └── 6. Return generated answer to user

```

## Azure Services Setup

### 1. Azure AI Search Index Schema

Create an index to store document chunks and their embeddings:

```
{
 "name": "sharepoint-docs",
 "fields": [
 {
 "name": "id",
 "type": "Edm.String",
 "key": true
 },
 {
 "name": "content",
 "type": "Edm.String",
 "searchable": true
 },
 {
 "name": "content_vector",
 "type": "Collection(Edm.Single)",
 "dimensions": 1536,
 "vectorSearchProfile": "default"
 },
 {
 "name": "source_file",
 "type": "Edm.String",
 "filterable": true,
 "facetable": true
 },
 {
 "name": "site_name",
 "type": "Edm.String",
 "filterable": true
 },
 {
 "name": "file_path",
 "type": "Edm.String"
 },
 {
 "name": "modified_date",
 "type": "Edm.DateTimeOffset",
 "filterable": true,
 "sortable": true
 },
 {
 "name": "chunk_index",
 "type": "Edm.Int32"
 }
 "vectorSearch": {
 "profiles": [
 {
 "name": "default",
 "algorithm": "hnsw-algorithm"
 }
]
 }
}
```

```

],
 "algorithms": [
 {
 "name": "hnsw-algorithm",
 "kind": "hnsw",
 "hnswParameters": {
 "metric": "cosine",
 "m": 4,
 "efConstruction": 400,
 "efSearch": 500
 }
 }
]
 }
}

```

## 2. Azure OpenAI Deployment

Required deployments in Azure OpenAI:

| Model                  | Deployment Name | Purpose                 |
|------------------------|-----------------|-------------------------|
| text-embedding-ada-002 | embeddings      | Convert text to vectors |
| gpt-4                  | gpt4            | Generate answers        |

## 3. Azure Function App

- **Runtime:** Python 3.11
- **Plan:** Consumption (pay-per-use) or Premium (for VNet integration)
- **Functions:**
  - [ProcessDocument](#) - Triggered by Event Grid
  - [QueryRAG](#) - HTTP trigger for user queries
  - [SyncSharePoint](#) - Timer trigger for periodic full sync

## Code Examples

### Document Processing Function

```

import azure.functions as func
from azure.search.documents import SearchClient
from azure.identity import DefaultAzureCredential
from openai import AzureOpenAI
from msgraph import GraphServiceClient
import os

Initialize clients
credential = DefaultAzureCredential()

```

```
openai_client = AzureOpenAI(
 azure_endpoint=os.getenv("AZURE_OPENAI_ENDPOINT"),
 api_key=os.getenv("AZURE_OPENAI_KEY"),
 api_version="2024-02-15-preview"
)

search_client = SearchClient(
 endpoint=os.getenv("AZURE_SEARCH_ENDPOINT"),
 index_name="sharepoint-docs",
 credential=credential
)

def process_document(file_id: str, site_id: str, filename: str):
 """Process a SharePoint document and add to search index."""

 # 1. Download from SharePoint
 graph_client = GraphServiceClient(credential)
 content =
 graph_client.sites[site_id].drive.items[file_id].content.get()

 # 2. Parse document based on file type
 text = parse_document(content, filename)

 # 3. Chunk the text
 chunks = chunk_text(text, chunk_size=800, overlap=100)

 # 4. Generate embeddings and prepare documents
 documents = []
 for i, chunk in enumerate(chunks):
 # Get embedding from Azure OpenAI
 embedding_response = openai_client.embeddings.create(
 input=chunk,
 model="text-embedding-ada-002"
)
 embedding = embedding_response.data[0].embedding

 documents.append({
 "id": f"{file_id}_{i}",
 "content": chunk,
 "content_vector": embedding,
 "source_file": filename,
 "site_name": site_id,
 "chunk_index": i
 })

 # 5. Upload to Azure AI Search
 search_client.upload_documents(documents)

 return len(documents)

def chunk_text(text: str, chunk_size: int = 800, overlap: int = 100) ->
list:
 """Split text into overlapping chunks."""
```

```
chunks = []
start = 0

while start < len(text):
 end = start + chunk_size
 chunk = text[start:end]

 # Try to break at sentence boundary
 if end < len(text):
 last_period = chunk.rfind('.')
 if last_period > chunk_size * 0.5:
 chunk = chunk[:last_period + 1]
 end = start + last_period + 1

 chunks.append(chunk.strip())
 start = end - overlap

return chunks

def parse_document(content: bytes, filename: str) -> str:
 """Parse document content based on file type."""
 ext = filename.lower().split('.')[-1]

 if ext == 'pdf':
 import fitz
 doc = fitz.open(stream=content, filetype="pdf")
 return "\n".join([page.get_text() for page in doc])

 elif ext == 'docx':
 from docx import Document
 import io
 doc = Document(io.BytesIO(content))
 return "\n".join([p.text for p in doc.paragraphs])

 elif ext == 'xlsx':
 from openpyxl import load_workbook
 import io
 wb = load_workbook(io.BytesIO(content))
 text_parts = []
 for sheet in wb:
 for row in sheet.iter_rows(values_only=True):
 row_text = " | ".join([str(c) for c in row if c])
 if row_text:
 text_parts.append(row_text)
 return "\n".join(text_parts)

 elif ext in ['txt', 'md']:
 return content.decode('utf-8', errors='ignore')

 return ""
```

## RAG Query Function

```
import azure.functions as func
import json

def query_rag(question: str, site_filter: str = None) -> dict:
 """
 Process a user question using RAG.

 Args:
 question: User's natural language question
 site_filter: Optional SharePoint site to filter results

 Returns:
 Dictionary with answer and sources
 """

 # 1. Generate embedding for the question
 question_embedding = openai_client.embeddings.create(
 input=question,
 model="text-embedding-ada-002"
).data[0].embedding

 # 2. Search Azure AI Search with vector query
 search_results = search_client.search(
 search_text=question,
 vector_queries=[{
 "vector": question_embedding,
 "k_nearest_neighbors": 5,
 "fields": "content_vector"
 }],
 filter=f"site_name eq '{site_filter}'" if site_filter else None,
 select=["content", "source_file", "file_path"]
)

 # 3. Build context from search results
 results_list = list(search_results)

 if not results_list:
 return {
 "answer": "I couldn't find any relevant information in the documents.",
 "sources": []
 }

 context_parts = []
 sources = []

 for result in results_list:
 context_parts.append(
 f"From '{result['source_file']}':\n{result['content']}"
)
```

```
 if result['source_file'] not in sources:
 sources.append(result['source_file'])

 context = "\n\n---\n\n".join(context_parts)

 # 4. Generate answer using Azure OpenAI
 system_prompt = """You are a helpful assistant that answers questions
based on
the provided document context.

Rules:
- Only answer based on the provided context
- If the context doesn't contain the answer, say so
- Be concise and direct
- Cite which document the information came from"""

 response = openai_client.chat.completions.create(
 model="gpt-4",
 temperature=0,
 messages=[
 {"role": "system", "content": system_prompt},
 {"role": "user", "content": f"Context:\n{context}\n\nQuestion:
{question}"}
]
)

 return {
 "answer": response.choices[0].message.content,
 "sources": sources
 }

HTTP Trigger for Azure Function
def main(req: func.HttpRequest) -> func.HttpResponse:
 try:
 body = req.get_json()
 question = body.get('question')
 site_filter = body.get('site')

 if not question:
 return func.HttpResponse(
 json.dumps({"error": "Question is required"}),
 status_code=400
)

 result = query_rag(question, site_filter)

 return func.HttpResponse(
 json.dumps(result),
 mimetype="application/json"
)

 except Exception as e:
 return func.HttpResponse(
```

```

 json.dumps({"error": str(e)}),
 status_code=500
)

```

## Cost Estimates

### Monthly Cost Breakdown

| Service                   | Tier                     | Estimated Cost                    |
|---------------------------|--------------------------|-----------------------------------|
| Azure AI Search           | Basic (15GB, 3 replicas) | \$75/month                        |
| Azure OpenAI - Embeddings | text-embedding-ada-002   | \$0.0001/1K tokens                |
| Azure OpenAI - GPT-4      | gpt-4                    | \$0.03/1K input, \$0.06/1K output |
| Azure Functions           | Consumption plan         | \$0-20/month                      |
| Azure Storage             | Standard                 | \$5/month                         |
| Event Grid                | Per operation            | ~\$1/month                        |

### Cost Scenarios

| Usage Level | Queries/Month | Documents | Est. Monthly Cost |
|-------------|---------------|-----------|-------------------|
| Light       | 1,000         | 500       | ~\$100            |
| Moderate    | 10,000        | 2,000     | ~\$150            |
| Heavy       | 50,000        | 10,000    | ~\$300            |

## Comparison with M365 Copilot

| Aspect                    | M365 Copilot      | Custom Azure RAG        |
|---------------------------|-------------------|-------------------------|
| <b>Cost (100 users)</b>   | \$3,000/month     | ~\$150/month            |
| <b>Cost (500 users)</b>   | \$15,000/month    | ~\$200/month            |
| <b>Setup Time</b>         | Immediate         | 2-4 weeks               |
| <b>Maintenance</b>        | Microsoft managed | Self-managed            |
| <b>Customization</b>      | Limited           | Full control            |
| <b>Prompt Engineering</b> | Not possible      | Fully customizable      |
| <b>Data Location</b>      | Microsoft cloud   | Your Azure tenant       |
| <b>Supported Sources</b>  | All M365 apps     | SharePoint (extensible) |

Annual Savings (100 users)

|                   |                              |
|-------------------|------------------------------|
| M365 Copilot:     | \$3,000 × 12 = \$36,000/year |
| Custom Azure RAG: | \$150 × 12 = \$1,800/year    |

---

|          |               |
|----------|---------------|
| Savings: | \$34,200/year |
|----------|---------------|

## UI Options

### Option 1: Power Apps (Recommended for Quick Start)

- **Effort:** Low (1-2 days)
- **Skills:** No coding required
- **Best for:** Internal business users
- **Features:** Forms, basic chat interface, SharePoint integration

### Option 2: Teams Bot

- **Effort:** Medium (1 week)
- **Skills:** Bot Framework, Node.js/C#
- **Best for:** Teams-centric organizations
- **Features:** Conversational UI, embedded in Teams

### Option 3: SharePoint Web Part (SPFx)

- **Effort:** Medium (1-2 weeks)
- **Skills:** React, TypeScript, SPFx
- **Best for:** Embedded experience in SharePoint
- **Features:** Native SharePoint look and feel

### Option 4: Custom Web Application

- **Effort:** High (2-4 weeks)
- **Skills:** Full-stack development
- **Best for:** Public-facing or highly custom needs
- **Features:** Complete flexibility

---

## Implementation Roadmap

### Phase 1: Foundation (Week 1)

- Set up Azure resource group
- Deploy Azure AI Search
- Deploy Azure OpenAI
- Configure Azure AD app registration
- Set up SharePoint API permissions

### Phase 2: Ingestion Pipeline (Week 2)

- Create Azure Function App
- Implement document processing function
- Set up Event Grid subscription for SharePoint
- Test with sample documents
- Initial document sync

### Phase 3: Query API (Week 3)

- Implement RAG query function
- Add authentication/authorization
- Implement rate limiting
- Test query accuracy
- Tune search parameters

### Phase 4: User Interface (Week 4)

- Build Power App / Teams Bot / Web UI
- Connect to RAG API
- Add source citations
- User acceptance testing
- Documentation and training

### Phase 5: Production (Ongoing)

- Monitor performance and costs
  - Tune prompts based on feedback
  - Add additional SharePoint sites
  - Implement feedback loop for improvement
- 

## Security Considerations

1. **Authentication:** Use Azure AD for all service-to-service auth
  2. **Authorization:** Respect SharePoint permissions in search results
  3. **Data Encryption:** Enable encryption at rest and in transit
  4. **Network:** Consider Private Endpoints for sensitive data
  5. **Logging:** Enable Azure Monitor for audit trails
  6. **API Security:** Use API Management for rate limiting and key management
- 

## Support and Resources

- [Azure AI Search Documentation](#)
  - [Azure OpenAI Documentation](#)
  - [Microsoft Graph API for SharePoint](#)
  - [Azure Functions Documentation](#)
-